

# **Reinforcement Learning:**

## **Basic concepts**

---


**Joelle Pineau**

School of Computer Science, McGill University  
Facebook AI Research (FAIR)


CIFAR Reinforcement Learning Summer School  
July 3 2017

# Reinforcement learning

- Learning by trial-and-error, in real-time.
- Improves with experience
- Inspired by psychology
  - Agent + Environment
  - Agent selects actions to maximize utility function.



# RL system circa 1990's: TD-Gammon



Reward function:

- +100 if win
- 100 if lose
- 0 for all other states

Trained by playing  $1.5 \times 10^6$  million games against itself.


Enough to beat the best human player.

# 2016: World Go Champion Beaten by Deep Learning



# RL applications at RLDM 2017

- Robotics
- Video games
- Conversational systems
- Medical intervention
- Algorithm improvement
- Improvisational theatre
- Autonomous driving
- Prosthetic arm control
- Financial trading
- Query completion




# When to use RL?


- Data in the form of trajectories.
- Need to make a sequence of (related) decisions.
- Observe (partial, noisy) feedback to choice of actions.
- Tasks that require both learning and planning.

# RL vs supervised learning

Training signal = desired (target outputs), e.g. class





Training signal = “rewards”





# RL vs supervised learning

Training signal = desired (target outputs), e.g. class



# RL vs supervised learning

Training signal = desired (target outputs), e.g. class



Practical & technical challenges:

1. Need access to the environment.
2. Jointly learning AND planning from **correlated** samples.
3. Data distribution changes with action choice.

# Markov Decision Process (MDP)

Defined by:


$S := \{s_1, s_2, \dots, s_n\}$ , the set of states (*can be infinite/continuous*)

$A := \{a_1, a_2, \dots, a_m\}$ , the set of actions (*can be infinite/continuous*)

$T(s,a,s') := Pr(s'|s,a)$ , the dynamics of the environment

$R(s,a)$ : Reward function


$\mu(s)$  : Initial state distribution



# The Markov property

The distribution over future states **depends only on the present state and action**, not on any other previous event.

$$Pr(s_{t+1} | s_0, \dots, s_t, a_0, \dots, a_t) = Pr(s_{t+1} | s_t, a_t)$$



# The **Markov** property

- Traffic lights?




- Chess?



# The Markov property

- Traffic lights?



- Chess?



- Poker?



**Tip:** Incorporate past observations in the state to have sufficient information to predict next state.

# The goal of RL? Maximize return!

- Return,  $U_t$  of a trajectory, is the sum of rewards starting from step  $t$ .

# The goal of RL? Maximize return!

- Return,  $U_t$  of a trajectory, is the sum of rewards starting from step  $t$ .
- **Episodic task:** consider return over finite horizon (e.g. games, maze).

$$U_t = r_t + r_{t+1} + r_{t+2} + \dots + r_T$$

- **Continuing task:** consider return over infinite horizon (e.g. juggling, balancing).

$$U_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} \dots = \sum_{k=0: \infty} \gamma^k r_{t+k}$$

# The discount factor, $\gamma$

- Discount factor,  $\gamma \in [0, 1)$  (usually close to 1).
- Intuition:
  - Receiving \$80 today is worth the same as \$100 tomorrow (assuming a discount factor of factor of  $\gamma = 0.8$ ).
  - At each time step, there is a  $1 - \gamma$  chance that the agent dies, and does not receive rewards afterwards.

# Defining behavior: The policy

- Policy,  $\pi$  defines the action-selection strategy at every state:


$$\pi(s, a) = P(a_t=a \mid s_t=s)$$

$$\pi : S \rightarrow A$$

Goal: **Find the policy that maximizes expected total reward.**  
*(But there are many policies!)*

$$\operatorname{argmax}_\pi E_\pi [ r_0 + r_1 + \dots + r_T \mid s_0 ]$$


# Example: Career Options



$n$ =Do Nothing  
 $i$  = Apply to industry  
 $g$  = Apply to grad school  
 $a$  = Apply to academia

What is the best policy?

# Example: Career Options



n=Do Nothing  
i = Apply to industry  
g = Apply to grad school  
a = Apply to academia

What is the best policy?

# Value functions

The **expected return of a policy** (for every state) is called the **value function**:  $V^\pi(s) = E_\pi [r_t + r_{t+1} + \dots + r_T | s_t = s]$

Simple strategy to find the best policy:

1. Enumerate the space of all possible policies.
2. Estimate the expected return of each one.
3. Keep the policy that has maximum expected return.

# Getting confused with terminology?

- **Reward?**
- **Return?**
- **Value?**
- **Utility?**

# Getting confused with terminology?

- **Reward**: 1 step numerical feedback
- **Return**: Sum of rewards over the agent's trajectory.
- **Value**: Expected sum of rewards over the agent's trajectory.
- **Utility**: Numerical function representing preferences.
- In RL, we assume **Utility = Return**.

# The value of a policy

$$V^\pi(s) = E_\pi [r_t + r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = E_\pi [r_t] + E_\pi [r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = \underbrace{\sum_{a \in A} \pi(s, a) R(s, a)}_{\text{Immediate reward}} + \underbrace{E_\pi [r_{t+1} + \dots + r_T \mid s_t = s]}_{\text{Future expected sum of rewards}}$$

# The value of a policy

$$V^\pi(s) = E_\pi [r_t + r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = E_\pi [r_t] + E_\pi [r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = \sum_{a \in A} \pi(s, a) R(s, a) + E_\pi [r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = \sum_{a \in A} \pi(s, a) R(s, a) + \underbrace{\sum_{a \in A} \pi(s, a) \sum_{s' \in S} T(s, a, s') E_\pi [r_{t+1} + \dots + r_T \mid s_{t+1} = s']}_{\text{Expectation over 1-step transition}}$$

# The value of a policy

$$V^\pi(s) = E_\pi [r_t + r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = E_\pi [r_t] + E_\pi [r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = \sum_{a \in A} \pi(s, a) R(s, a) + E_\pi [r_{t+1} + \dots + r_T \mid s_t = s]$$

$$V^\pi(s) = \sum_{a \in A} \pi(s, a) R(s, a) + \sum_{a \in A} \pi(s, a) \sum_{s' \in S} T(s, a, s') \underbrace{E_\pi [r_{t+1} + \dots + r_T \mid s_{t+1} = s']}$$

$$V^\pi(s) = \sum_{a \in A} \pi(s, a) R(s, a) + \sum_{a \in A} \pi(s, a) \sum_{s' \in S} T(s, a, s') V^\pi(s')$$

*By definition*

This is a **dynamic programming** algorithm.

# The value of a policy

State value function (for a **fixed** policy):

$$V^\pi(s) = \sum_{a \in A} \pi(s, a) [ R(s, a) + \gamma \underbrace{\sum_{s' \in S} T(s, a, s') V^\pi(s')}_{\text{Immediate Future expected sum of rewards}} ]$$

State-action value function:

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') [\sum_{a' \in A} \pi(s', a') Q^\pi(s', a')]$$

These are two forms of **Bellman's equation**.

# The value of a policy

State value function:

$$V^\pi(s) = \sum_{a \in A} \pi(s, a) (R(s, a) + \gamma \sum_{s' \in S} T(s, a, s') V^\pi(s'))$$

When  $S$  is a **finite set of states**, this is a **system of linear equations** (one per state) with a unique solution  $V^\pi$ .

Bellman's equation in matrix form:

$$V^\pi = R^\pi + \gamma T^\pi V^\pi$$

Which can solved exactly:

$$V^\pi = (I - \gamma T^\pi)^{-1} R^\pi$$

# Iterative Policy Evaluation: Fixed policy

Main idea: turn Bellman equations into update rules.

1. Start with some initial guess  $V_0(s)$ ,  $\forall s$ . (Can be 0, or  $r(s, \cdot)$ .)

# Iterative Policy Evaluation: Fixed policy

Main idea: turn Bellman equations into update rules.

1. Start with some initial guess  $V_0(s)$ ,  $\forall s$ . (Can be 0, or  $r(s, \cdot)$ .)
2. During every iteration  $k$ , update the value function for all states:

$$V_{k+1}(s) \leftarrow \left( R(s, \pi(s)) + \gamma \sum_{s' \in S} T(s, \pi(s), s') V_k(s') \right)$$

# Iterative Policy Evaluation: Fixed policy

Main idea: turn Bellman equations into update rules.

1. Start with some initial guess  $V_0(s)$ ,  $\forall s$ . (Can be 0, or  $r(s, \cdot)$ .)
2. During every iteration  $k$ , update the value function for all states:

$$V_{k+1}(s) \leftarrow \left( R(s, \pi(s)) + \gamma \sum_{s' \in S} T(s, \pi(s), s') V_k(s') \right)$$

3. Stop when the maximum changes between two iterations is smaller than a desired threshold (the values stop changing.)

**This is a dynamic programming algorithm. Guaranteed to converge!**

# Convergence of Iterative Policy Evaluation

- Consider the absolute error in our estimate  $V_{k+1}(s)$ :

$$\begin{aligned}|V_{k+1}(s) - V^\pi(s)| &= \left| \sum_a \pi(s, a)(R(s, a) + \gamma \sum_{s'} T(s, a, s') V_k(s')) \right. \\&\quad \left. - \sum_a \pi(s, a)(R(s, a) + \gamma \sum_{s'} T(s, a, s') V^\pi(s')) \right| \\&= \gamma \left| \sum_a \pi(s, a) \sum_{s'} T(s, a, s') (V_k(s') - V^\pi(s')) \right| \\&\leq \gamma \sum_a \pi(s, a) \sum_{s'} T(s, a, s') |V_k(s') - V^\pi(s')|\end{aligned}$$

- As long as  $\gamma < 1$ , the error **contracts** and eventually goes to 0.

# Optimal policies and optimal value functions

- **Optimal value function**,  $V^*$  is the highest value that can be achieved for each state:

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

- Any policy that achieves  $V^*$  is called an **optimal policy**,  $\pi^*$ .

# Optimal policies and optimal value functions

- **Optimal value function**,  $V^*$  is the highest value that can be achieved for each state:

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

- Any policy that achieves  $V^*$  is called an **optimal policy**,  $\pi^*$ .
- For each MDP there is a **unique optimal value function** (*Bellman, 1957*).
- The optimal policy is not necessarily unique.

# Optimal policies and optimal value functions

- If we know  $V^*$  (and  $R, T, \gamma$ ), then we can compute  $\pi^*$  easily.

$$\pi^*(s) = \operatorname{argmax}_{a \in A} (R(s,a) + \gamma \sum_{s' \in S} T(s,a,s')V^*(s'))$$

- If we know  $\pi^*$  (and  $R, T, \gamma$ ), then we can compute  $V^*$  easily.

$$V^*(s) = \sum_{a \in A} \pi^*(s,a) (R(s,a) + \gamma \sum_{s' \in S} T(s,a,s')V^*(s'))$$

$$V^*(s) = R(s, \pi(s)) + \gamma \sum_{s' \in S} T(s, \pi(s), s')V^*(s')$$

**Take-home:** Both  $V^*$  and  $\pi^*$  are “solutions” to the MDP.

# Finding a good policy: Policy Iteration

- Start with an initial policy  $\pi_0$  (e.g. random)
- Repeat:
  - Compute  $V^\pi$ , using iterative policy evaluation.
  - Compute a new policy  $\pi'$  that is greedy with respect to  $V^\pi$
- Terminate when  $\pi = \pi'$

# Finding a good policy: **Value iteration**

Main idea: Turn the Bellman optimality equation into an iterative update rule (same as done in policy evaluation):

1. Start with an arbitrary initial approximation  $V_0(s)$
2. On each iteration, update the value function estimate:  
$$V_k(s) = \max_{a \in A} ( R(s,a) + \gamma \sum_{s' \in S} T(s,a,s') V_{k-1}(s') )$$
3. Stop when max value change between iterations is below threshold.

The algorithm converges (in the limit) to the true  $V^*$ .

# Three related algorithms


1. **Policy evaluation:** Fix the policy, estimate its value.
2. **Policy iteration:** Find the best policy at each state.
  - » Policy evaluation + greedy improvement.
3. **Value iteration:** Find the optimal value function.

# Three related algorithms

1. **Policy evaluation:** Fix the policy, estimate its value.
  - $O(S^3)$
2. **Policy iteration:** Find the best policy at each state.
  - » Policy evaluation + greedy improvement.
  - $O(S^3+S^2A)$  per iteration
3. **Value iteration:** Find the optimal value function.
  - $O(S^2A)$  per iteration

# A 4x3 gridworld example

- 11 discrete states, 4 motion actions (N, S, E, W) in each state.
- Transitions are mildly **stochastic**.
- Reward is +1 in top right state, -10 in state directly below, -0 elsewhere.
- Episode terminates when the agent reaches +1 or -10 state.
- Discount factor  $\gamma = 0.99$ .



S			+1
			-10

# Value Iteration (1)

0	0	0	+1
0		0	-10
0	0	0	0

## Value Iteration (2)

0	0	0.69	+1
0		-0.99	-10
0	0	0	-0.99

Bellman residual:  $|V_2(s) - V_1(s)| = 0.99$

# Value Iteration (5)

0.48	0.70	0.76	+1
0.23		-0.55	-10
0	-0.20	-0.23	-1.40

Bellman residual:  $|V_5(s) - V_4(s)| = 0.23$


# Value Iteration (20)

0.78	0.80	0.81	+1
0.77		-0.44	-10
0.75	0.69	0.37	-0.92


Bellman residual:  $|V_5(s) - V_4(s)| = 0.008$

# Another example: Four Rooms


- Four actions, fail 30% of the time.
- No rewards until the goal is reached,  $\gamma = 0.9$ .
- Values propagate backwards from the goal.



Iteration #1



Iteration #2





Iteration #3

# Asynchronous value iteration

- Instead of updating all states on every iteration, focus on *important states*.
  - E.g., board positions that occur on every game, rather than just once in 100 games.
- Asynchronous dynamic programming algorithm:
  - Generate trajectories through the MDP.
  - Update states whenever they appear on such a trajectory.
- Focuses the updates on states that are actually possible.

# Generalized Policy Iteration

- Any combination of policy evaluation and policy improvement steps.  
e.g. only update value of one state and improve policy at that state.




# Key challenges in RL

- Designing the problem domain
  - State representation
  - Action choice
  - Cost/reward signal
- Acquiring data for training
  - Exploration / exploitation
  - High cost actions
  - Time-delayed cost/reward signal
- Function approximation
- Validation / confidence measures



# Learning online from trial & error



# Online reinforcement learning

- **Monte-Carlo** value estimate: Use the empirical return,  $U(s_t)$  as a target estimate for the actual value function:

$$V(s_t) \leftarrow V(s_t) + \alpha (U(s_t) - V(s_t))$$

*\* Not a Bellman equation. More like a gradient equation.*


- Here  $\alpha$  is the learning rate (a parameter).
- Need to wait until the end of the trajectory to compute  $U(s_t)$ .

# Temporal-Difference (TD) learning


- Monte-Carlo learning:  $V(s_t) \leftarrow V(s_t) + \alpha(U(s_t) - V(s_t))$

- TD-learning:

$$V(s_t) \leftarrow V(s_t) + \alpha(r_{t+1} + \gamma V(s_{t+1}) - V(s_t)) \quad \forall t = 0, 1, 2, \dots$$



# TD-Gammon (Tesauro, 1992)



Reward function:

- +100 if win
- 100 if lose
- 0 for all other states

Trained by playing  $1.5 \times 10^6$  million games against itself.

Enough to beat the best human player.



# Several challenges in RL

- Designing the problem domain
  - State representation
  - Action choice
  - Cost/reward signal
- Acquiring data for training
  - Exploration / exploitation
  - High cost actions
- Time-delayed cost/reward signal
- **Function approximation**
- Validation / confidence measures



# Tabular / Function approximation

- **Tabular:** Can store in memory a list of the states and their value.



\* *Can prove many more theoretical properties in this case, about convergence, sample complexity.*

- **Function approximation:** Too many states, continuous state spaces.



In large state spaces: Need approximation


Challenge: finding good features

$$\hat{Q}^\pi(s, a) = \sum_{i=1}^d \theta_i \underline{\phi_i(s, a)}$$


feature vector

# Learning representations for RL

$s$



Original state




$$Q_{\theta}(s, a)$$




Linear function

# Deep Reinforcement Learning

$S$



Original state





Convolutional Neural Net

Deep Q-Network trained with stochastic gradient descent.

[DeepMind: Mnih et al., 2015].

# Deep RL in Minecraft



Many possible architectures,  
incl. **memory** and **context**

Online videos: <https://sites.google.com/a/umich.edu/junhyuk-oh/icml2016-minecraft>

[U.Michigan: Oh et al., 2016].

# The RL lingo

- Episodic / Continuing task
- Batch / Online
- **On-policy / Off-policy**
- Exploration / Exploitation
- Model-based / Model-free
- Policy optimization / Value function methods

# On-policy / Off-policy

- Policy induces a distribution over the states (data).
  - Data distribution **changes** every time you change the policy!

# On-policy / Off-policy



- Policy induces a distribution over the states (data).
  - Data distribution **changes** every time you change the policy!
- Evaluating several policies with the same batch:
  - Need very big batch!
  - Need policy to adequately cover all  $(s,a)$  pairs.

# On-policy / Off-policy


- Policy induces a distribution over the states (data).
  - Data distribution **changes** every time you change the policy!
- Evaluating several policies with the same batch:
  - Need very big batch!
  - Need policy to adequately cover all  $(s,a)$  pairs.
- Use importance sampling to reweigh data samples to compute unbiased estimates of a new policy.

$$\rho_t = \frac{\pi(s_t, a_t)}{b(s_t, a_t)}$$

# Exploration / Exploitation




# Exploration / Exploitation



# Model-based vs Model-free RL


- **Option #1:** Collect large amounts of observed trajectories.  
Learn an approximate model of the dynamics (e.g. with supervised learning). Pretend the model is correct and apply value iteration.
- **Option #2:** Use data to directly learn the value function or optimal policy.

# Policy Optimization / Value Function



# Quick summary

- RL problems are everywhere!
  - Games, text, robotics, medicine, ...
- Need access to the “environment” to generate samples.
  - Most recent results make extensive use of a simulator.
- Feasible methods for large, complex tasks.
- Intuition about what is “easy”, “hard” is different than supervised learning.



# RL resources

Comprehensive list of resources:

- <https://github.com/aikorea/awesome-rl>

Environments & algorithms:

- [http://glue.rl-community.org/wiki/Main\\_Page](http://glue.rl-community.org/wiki/Main_Page)
- <https://gym.openai.com>
- <https://github.com/deepmind/lab>

