

Learning Structured Sparsity in Deep Neural Networks

Wei Wen¹ Chunpeng Wu¹ Yandan Wang¹ Yiran Chen¹ Hai
Li¹

¹University of Pittsburgh

NIPS, 2016

Presenter: Bargav Jayaraman

Outline

- 1 Introduction
- 2 Related Works
- 3 Proposed Structure Sparsity Learning Approach
 - SSL for Generic Structures
 - SSL for Filters and Channels
 - SSL for Filter Shapes
 - SSL for Layer Depth
 - SSL for Computationally Efficient Structures
- 4 Experimental Results
- 5 Summary

Introduction

- **Problem:** Deployment of large-scale deep learning model is computationally expensive
- **Solution:** Occam's Razor - Simple is better!
Remove or zero-out the non-essential weights / layers of the model

Introduction

- **Problem:** Deployment of large-scale deep learning model is computationally expensive
- **Solution:** Occam's Razor - Simple is better!
Remove or zero-out the non-essential weights / layers of the model
Catch: Trade-off between model complexity and accuracy

Related Works

- **Connection pruning and weight sparsifying.** Connection pruning removes unwanted weight connections from the fully connected layers of a CNN. **Not much beneficial for convolutional layers!** Hard-coding sparse weights for convolutional layers introduces non-structured sparsity with slight accuracy loss.
 - This work achieves structured sparsity in adjacent memory space

- **Connection pruning and weight sparsifying.** Connection pruning removes unwanted weight connections from the fully connected layers of a CNN. **Not much beneficial for convolutional layers!**
Hard-coding sparse weights for convolutional layers introduces non-structured sparsity with slight accuracy loss.
 - This work achieves structured sparsity in adjacent memory space
- **Low rank approximation.** LRA compresses the deep network by decomposing the weight matrix $W \in \mathbb{R}^{u \times v}$ at every layer into product of two matrices $U \in \mathbb{R}^{u \times \alpha}$ and $V \in \mathbb{R}^{\alpha \times v}$, where $\alpha < u, v$.
 - This work dynamically optimizes the model and obtains lower rank approximation

- **Connection pruning and weight sparsifying.** Connection pruning removes unwanted weight connections from the fully connected layers of a CNN. **Not much beneficial for convolutional layers!**
Hard-coding sparse weights for convolutional layers introduces non-structured sparsity with slight accuracy loss.
 - This work achieves structured sparsity in adjacent memory space
- **Low rank approximation.** LRA compresses the deep network by decomposing the weight matrix $W \in \mathbb{R}^{u \times v}$ at every layer into product of two matrices $U \in \mathbb{R}^{u \times \alpha}$ and $V \in \mathbb{R}^{\alpha \times v}$, where $\alpha < u, v$.
 - This work dynamically optimizes the model and obtains lower rank approximation
- **Model structure learning.** Group Lasso has been used for structure sparsity in deep models to learn the appropriate number of filters or filter shapes.
 - This work applies group Lasso at various levels of the deep model

Outline

- 1 Introduction
- 2 Related Works
- 3 **Proposed Structure Sparsity Learning Approach**
 - **SSL for Generic Structures**
 - SSL for Filters and Channels
 - SSL for Filter Shapes
 - SSL for Layer Depth
 - SSL for Computationally Efficient Structures
- 4 Experimental Results
- 5 Summary

Structure Sparsity Learning for Generic Structures

Consider the weights of a deep network as a 4-D tensor:

$W^{(l)} \in \mathbb{R}^{N_l \times C_l \times M_l \times K_l}$, where N_l , C_l , M_l and K_l are the dimensions of the l -th layer ($1 \leq l \leq L$) weight tensor along the axes of filter, channel, spatial height and spatial width. L denotes the number of convolutional layers. Then the proposed generic optimization is:

$$E(W) = E_D(W) + \lambda.R(W) + \lambda_g \cdot \sum_{l=1}^L R_g(W^{(l)})$$

$E_D(W)$ is the loss on data, $R(\cdot)$ is the non-structured regularizer, like l_2 -norm, and $R_g(\cdot)$ is the structured regularizer. This work uses group Lasso for $R_g(\cdot)$.

Group Lasso

- The regularization of group Lasso on a set of weights w is given as: $R_g(w) = \sum_{g=1}^G \|w^{(g)}\|_g$, where g is a group of partial weights in w and G is the total number of groups.
- $\|\cdot\|_g$ is the group Lasso, or $\|w^{(g)}\|_g = \sqrt{\sum_{i=1}^{|w^{(g)}|} (w_i^{(g)})^2}$, where $|w^{(g)}|$ is the number of weights in $w^{(g)}$.

- The regularization of group Lasso on a set of weights w is given as: $R_g(w) = \sum_{g=1}^G \|w^{(g)}\|_g$, where g is a group of partial weights in w and G is the total number of groups.
- $\|\cdot\|_g$ is the group Lasso, or $\|w^{(g)}\|_g = \sqrt{\sum_{i=1}^{|w^{(g)}|} (w_i^{(g)})^2}$, where $|w^{(g)}|$ is the number of weights in $w^{(g)}$.

Question: Why is this called group “Lasso” if it uses l_2 -regularization?

Group Lasso

- The regularization of group Lasso on a set of weights w is given as: $R_g(w) = \sum_{g=1}^G \|w^{(g)}\|_g$, where g is a group of partial weights in w and G is the total number of groups.
- $\|\cdot\|_g$ is the group Lasso, or $\|w^{(g)}\|_g = \sqrt{\sum_{i=1}^{|w^{(g)}|} (w_i^{(g)})^2}$, where $|w^{(g)}|$ is the number of weights in $w^{(g)}$.

Question: Why is this called group “Lasso” if it uses l_2 -regularization?

Answer: l_2 -regularization has all-or-none zero effect!


Outline

- 1 Introduction
- 2 Related Works
- 3 Proposed Structure Sparsity Learning Approach**
 - SSL for Generic Structures
 - SSL for Filters and Channels**
 - SSL for Filter Shapes
 - SSL for Layer Depth
 - SSL for Computationally Efficient Structures
- 4 Experimental Results
- 5 Summary

SSL for Filters and Channels

Suppose $W_{n_l, :, :, :}^{(l)}$ is the n_l -th filter and $W_{:, c_l, :, :}^{(l)}$ is the c_l -th channel of all filters in the l -th layer. Then the optimization target is defined as:

$$E(W) = E_D(W) + \lambda_n \cdot \sum_{l=1}^L \left(\sum_{n_l=1}^{N_l} \|W_{n_l, :, :, :}^{(l)}\|_g \right) + \lambda_c \cdot \sum_{l=1}^L \left(\sum_{c_l=1}^{C_l} \|W_{:, c_l, :, :}^{(l)}\|_g \right)$$




Outline

- 1 Introduction
- 2 Related Works
- 3 Proposed Structure Sparsity Learning Approach**
 - SSL for Generic Structures
 - SSL for Filters and Channels
 - SSL for Filter Shapes**
 - SSL for Layer Depth
 - SSL for Computationally Efficient Structures
- 4 Experimental Results
- 5 Summary

SSL for Filter Shapes

Suppose $W_{:,c_l,m_l,k_l}^{(l)}$ denotes the vector of all corresponding weights of spatial position (m_l, k_l) in the filters across c_l -th channel, then:

$$E(W) = E_D(W) + \lambda_s \cdot \sum_{l=1}^L \left(\sum_{c_l=1}^{C_l} \sum_{m_l=1}^{M_l} \sum_{k_l=1}^{K_l} \|W_{:,c_l,m_l,k_l}^{(l)}\|_g \right)$$



Outline


- 1 Introduction
- 2 Related Works
- 3 Proposed Structure Sparsity Learning Approach**
 - SSL for Generic Structures
 - SSL for Filters and Channels
 - SSL for Filter Shapes
 - SSL for Layer Depth**
 - SSL for Computationally Efficient Structures
- 4 Experimental Results
- 5 Summary

SSL for Layer Depth

Depth sparsity reduces the computation cost and improves accuracy. The optimization is given as:

$$E(W) = E_D(W) + \lambda_d \cdot \sum_{l=1}^L \|W^{(l)}\|_g$$

Zeroing out all filters in a layer can hinder the message passing across layers, and hence shortcut is used to transfer the feature map.



Outline

- 1 Introduction
- 2 Related Works
- 3 Proposed Structure Sparsity Learning Approach**
 - SSL for Generic Structures
 - SSL for Filters and Channels
 - SSL for Filter Shapes
 - SSL for Layer Depth
 - **SSL for Computationally Efficient Structures**
- 4 Experimental Results
- 5 Summary

SSL for Computationally Efficient Structures

- **2D-filter-wise sparsity for convolution.** Fine-grain variant of filter-wise sparsity is zeroing out 2D filters instead of 3D filters for efficient computation reduction. Since, 2D filters are smaller groups and hence easy to zero-out.
- **Combination of filter-wise and shape-wise sparsity for GEMM.** Convolutional operation is represented as a matrix in GEneral Matrix Multiplication (GEMM) such that each row is represented as a feature and each column is a collection of weight corresponding to shape sparsity. Combining filter-wise and shape-wise sparsity zeroes out the rows and columns of the weight matrix and hence reduces the dimensionality.

Experimental Results


- Filter-wise, Channel-wise and Shape-wise SSL on LeNet
- SSL on fully-connected MLP
- Filter-wise and Shape-wise SSL on ConvNet
- Depth-wise SSL on ResNet
- SSL on AlexNet

Table 1: Results after penalizing unimportant filters and channels in *LeNet*

| <i>LeNet</i> # | Error | Filter # [§] | Channel # [§] | FLOP [§] | Speedup [§] |
|-----------------------|-------|-----------------------|------------------------|-------------------|----------------------|
| 1 (<i>baseline</i>) | 0.9% | 20—50 | 1—20 | 100%—100% | 1.00×—1.00× |
| 2 | 0.8% | 5—19 | 1—4 | 25%—7.6% | 1.64×—5.23× |
| 3 | 1.0% | 3—12 | 1—3 | 15%—3.6% | 1.99×—7.44× |

[§]In the order of *conv1*—*conv2*Table 2: Results after learning filter shapes in *LeNet*


| <i>LeNet</i> # | Error | Filter size [§] | Channel # | FLOP | Speedup |
|-----------------------|-------|--------------------------|-----------|-----------|--------------|
| 1 (<i>baseline</i>) | 0.9% | 25—500 | 1—20 | 100%—100% | 1.00×—1.00× |
| 4 | 0.8% | 21—41 | 1—2 | 8.4%—8.2% | 2.33×—6.93× |
| 5 | 1.0% | 7—14 | 1—1 | 1.4%—2.8% | 5.19×—10.82× |

[§]The sizes of filters after removing zero shape fibers, in the order of *conv1*—*conv2*Figure 3: Learned *conv1* filters in *LeNet* 1 (top), *LeNet* 2 (middle) and *LeNet* 3 (bottom)

| <i>MLP #</i> | Error | Neuron # per layer [§] | FLOP per layer [§] |
|-----------------------|-------|---------------------------------|-----------------------------|
| 1 (<i>baseline</i>) | 1.43% | 784–500–300–10 | 100%–100%–100% |
| 2 | 1.34% | 469–294–166–10 | 35.18%–32.54%–55.33% |
| 3 | 1.53% | 434–174–78–10 | 19.26%–9.05%–26.00% |

[§]In the order of *input layer–hidden layer 1–hidden layer 2–output layer*

(a)



(b)

Figure 4: (a) Results of learning the number of neurons in *MLP*. (b) the connection numbers of input neurons (*i.e.* pixels) in *MLP 2* after SSL.

Table 3: Learning row-wise and column-wise sparsity of *ConvNet* on CIFAR-10

| <i>ConvNet</i> # | Error | Row sparsity [§] | Column sparsity [§] | Speedup [§] |
|-----------------------|-------|---------------------------|------------------------------|----------------------|
| 1 (<i>baseline</i>) | 17.9% | 12.5%–0%–0% | 0%–0%–0% | 1.00×–1.00×–1.00× |
| 2 | 17.9% | 50.0%–28.1%–1.6% | 0%–59.3%–35.1% | 1.43×–3.05×–1.57× |
| 3 | 16.9% | 31.3%–0%–1.6% | 0%–42.8%–9.8% | 1.25×–2.01×–1.18× |

[§]in the order of *conv1*–*conv2*–*conv3*

Figure 5: Learned *conv1* filters in *ConvNet* 1 (top), *ConvNet* 2 (middle) and *ConvNet* 3 (bottom)




Figure 6: Error vs. layer number after depth regularization by SSL. *ResNet-#* is the original *ResNet* in [5] with # layers. *SSL-ResNet-#* is the depth-regularized *ResNet* by SSL with # layers, including the last fully-connected layer. 32×32 indicates the convolutional layers with an output map size of 32×32 , and so forth.

Table 4: Sparsity and speedup of *AlexNet* on ILSVRC 2012

| # | Method | Top1 err. | Statistics | conv1 | conv2 | conv3 | conv4 | conv5 |
|---|------------|-----------|-----------------|-------|-------|-------|-------|-------|
| 1 | ℓ_1 | 44.67% | sparsity | 67.6% | 92.4% | 97.2% | 96.6% | 94.3% |
| | | | CPU \times | 0.80 | 2.91 | 4.84 | 3.83 | 2.76 |
| | | | GPU \times | 0.25 | 0.52 | 1.38 | 1.04 | 1.36 |
| 2 | SSL | 44.66% | column sparsity | 0.0% | 63.2% | 76.9% | 84.7% | 80.7% |
| | | | row sparsity | 9.4% | 12.9% | 40.6% | 46.9% | 0.0% |
| | | | CPU \times | 1.05 | 3.37 | 6.27 | 9.73 | 4.93 |
| | | | GPU \times | 1.00 | 2.37 | 4.94 | 4.03 | 3.05 |
| 3 | pruning[7] | 42.80% | sparsity | 16.0% | 62.0% | 65.0% | 63.0% | 63.0% |
| 4 | ℓ_1 | 42.51% | sparsity | 14.7% | 76.2% | 85.3% | 81.5% | 76.3% |
| | | | CPU \times | 0.34 | 0.99 | 1.30 | 1.10 | 0.93 |
| | | | GPU \times | 0.08 | 0.17 | 0.42 | 0.30 | 0.32 |
| 5 | SSL | 42.53% | column sparsity | 0.00% | 20.9% | 39.7% | 39.7% | 24.6% |
| | | | CPU \times | 1.00 | 1.27 | 1.64 | 1.68 | 1.32 |
| | | | GPU \times | 1.00 | 1.25 | 1.63 | 1.72 | 1.36 |

Summary

- Filter-wise, channel-wise, shape-wise and depth-wise SSL
- Dynamic compact structure learning without loss of accuracy
- Significant speed-ups with both CPUs and GPUs